

# Mathematical Models for AI News Bot

Dmitrii Bogolaev, Mark Lyumanov

May 2025

## 1 Mathematical Models

### 1.1 TF-IDF + KMeans Clustering

This model combines two powerful tools from natural language processing and unsupervised learning:

- **TF-IDF (Term Frequency–Inverse Document Frequency)** — transforms raw text into weighted vectors that emphasize rare but informative terms.
- **KMeans Clustering** — groups these vectors into clusters representing different topics.

The goal is to automatically select a small, diverse subset of articles that represent the main themes present in a larger set.

#### TF-IDF Explanation

TF-IDF evaluates how important a word is to a document in a collection. It is a product of two terms:

##### Term Frequency (TF):

$$\text{TF}(t, d) = \frac{f_{t,d}}{\sum_k f_{k,d}}$$

Here,  $f_{t,d}$  is the frequency of term  $t$  in document  $d$ , and the denominator sums frequencies of all terms in the same document.

##### Inverse Document Frequency (IDF):

$$\text{IDF}(t) = \log \left( \frac{N}{1 + n_t} \right)$$

Where  $N$  is the total number of documents and  $n_t$  is the number of documents in which term  $t$  appears. This gives higher weight to terms that are rare across the corpus.

##### TF-IDF Score:

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \cdot \text{IDF}(t)$$

The result is a high-dimensional vector for each document, with each component representing the TF-IDF weight of a term.

#### KMeans Clustering

KMeans is an iterative algorithm that aims to partition the dataset into  $k$  clusters by minimizing the variance within each cluster.

Given a set of vectors  $x_1, x_2, \dots, x_n$ , the goal is to find cluster centroids  $\mu_1, \mu_2, \dots, \mu_k$  such that:

$$\arg \min_{\{\mu\}} \sum_{i=1}^n \|x_i - \mu_{C_i}\|^2$$

Where  $C_i$  denotes the cluster assignment of  $x_i$ .

### Steps:

1. Initialize  $k$  random centroids.
2. Assign each point  $x_i$  to the nearest centroid.
3. Recompute each centroid as the mean of all points assigned to it.
4. Repeat steps 2 and 3 until convergence.

### Implementation Code

```
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.cluster import KMeans

def select_important(articles, n_clusters=3):
    if len(articles) == 0:
        return []

    texts = [article.get('title', '') for article in articles]

    if len(articles) < n_clusters:
        return articles

    vectorizer = TfidfVectorizer(stop_words='english')
    X = vectorizer.fit_transform(texts)

    kmeans = KMeans(n_clusters=n_clusters, random_state=42, n_init='auto')
    labels = kmeans.fit_predict(X)

    selected = []
    seen_labels = set()

    for i, label in enumerate(labels):
        if label not in seen_labels:
            selected.append(articles[i])
            seen_labels.add(label)

    return selected
```

### Why This Works

Articles that talk about similar topics tend to share key terms. TF-IDF helps identify these key terms, while KMeans groups together articles that emphasize the same vocabulary. By selecting one article per cluster, we create a concise and non-redundant summary of the broader information landscape.

This technique ensures topic diversity and avoids echoing similar headlines, which is crucial for a Telegram bot designed to inform rather than overwhelm.

## 1.2 Z-Score Statistical Selection

This model applies classical statistical hypothesis testing to detect news articles that significantly stand out in terms of their frequency or mention count.

### Motivation

When a particular topic is mentioned far more frequently than others, it may indicate a trending or important story. To quantify how unusual a value is, we use the **z-score**.

## Z-Score Formula

The z-score of an observation  $x$  is given by:

$$z = \frac{x - \mu}{\sigma}$$

Where:

- $x$  is the observed value (e.g., article mention count),
- $\mu$  is the mean of all values,
- $\sigma$  is the standard deviation of values.

A z-score tells us how many standard deviations  $x$  is away from the mean.

**Significance Threshold** For a typical 95% confidence level, we consider an observation significant if:

$$z > 1.96$$

This means the value is in the top 2.5% of the normal distribution — statistically rare under the assumption of normality.

## Practical Use Case

In our case, each article has a field `mention_count` (e.g., how many different sources mention it). We use z-scores to detect articles with unusually high mention counts and select them for summarization.

## Implementation Code

```
import numpy as np

def compute_z_score(x: float, mean: float, std: float) -> float:
    """
    Computes the z-score for a given value.
    """
    if std == 0:
        return 0.0
    return (x - mean) / std

def select_important(articles, threshold: float = 1.96):
    """
    Filters articles whose mention counts are significantly above the mean.
    """
    mention_counts = [a.get("mention_count", 1) for a in articles]

    if len(mention_counts) < 2:
        return articles

    mean = np.mean(mention_counts)
    std = np.std(mention_counts)

    filtered = []
    for article in articles:
        count = article.get("mention_count", 1)
        z = compute_z_score(count, mean, std)
        if z > threshold:
            filtered.append(article)

    return filtered
```

## Why This Works

This method is simple, interpretable, and grounded in statistical theory. It allows the bot to focus on content that exhibits strong anomalies in attention — likely reflecting emerging or viral topics — without needing any labeled data or training.

It is especially useful when your system monitors hundreds of articles and needs a principled way to filter out noise.

## 1.3 Bayesian Importance Estimation

This model estimates the probability that an article is important given its observed characteristics, using Bayes' Theorem from probability theory.

### Bayes' Theorem

Given a hypothesis  $H$  (the article is important) and observed data  $D$  (mention count), Bayes' rule defines:

$$P(H | D) = \frac{P(D | H) \cdot P(H)}{P(D)}$$

Where:

- $P(H)$  is the prior probability that a randomly selected article is important,
- $P(D | H)$  is the likelihood of observing the mention count if the article is important,
- $P(D)$  is the marginal probability of observing such data across all articles,
- $P(H | D)$  is the posterior — our updated belief in the article's importance.

### Implementation Assumptions

For simplicity:

- The prior is fixed at  $P(H) = 0.5$ ,
- The likelihood is approximated as a normalized count:  $P(D | H) \approx \frac{\text{mention\_count}}{10}$ ,
- The evidence is fixed as  $P(D) = 1$ , acting as a scaling constant.

## Python Code

```
def bayesian_importance(prior: float, likelihood: float, evidence: float) -> float:
    """
    Applies Bayes' theorem:
     $P(H/D) = (P(D/H) * P(H)) / P(D)$ 
    """
    if evidence == 0:
        return 0.0
    return (likelihood * prior) / evidence

def select_important(articles, threshold: float = 0.6):
    """
    Selects articles with posterior importance probability > threshold.
    """
    prior = 0.5          # Prior belief: 50% chance an article is important
    evidence = 1.0        # Simplified: fixed marginal likelihood

    selected = []
    for article in articles:
        mention_count = article.get("mention_count", 1)
        likelihood = min(mention_count / 10, 1.0) # Normalize to [0, 1]
        posterior = bayesian_importance(prior, likelihood, evidence)
        if posterior > threshold:
            selected.append(article)

    return selected
```

## Why This Works

This method formalizes a probabilistic belief update: the more often an article is mentioned, the higher the likelihood it is important. Bayes' Theorem combines this with a prior belief to yield a robust, interpretable metric.

It is particularly useful when explicit thresholds or supervised labels are not available but some notion of “probable importance” is still desired.

## 1.4 Logistic Regression Classification

This model uses supervised learning to predict whether a news article is important based on its numerical features. The method is based on **logistic regression**, a statistical model used for binary classification.

### Mathematical Foundation

The logistic regression hypothesis function is defined as:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Where:

- $x$  is the feature vector (e.g., title length, mention count),
- $\theta$  is the parameter vector (learned from data),
- $h_{\theta}(x) \in (0, 1)$  is the predicted probability that the article is important.

A threshold (e.g. 0.7) is used to convert the probability into a binary decision.

## Training Data

To train the model, we need a set of labeled data:

- Each sample represents an article with features such as:
  - Title length
  - Mention count
  - Score from platform (optional)
  - Keyword presence
- Each sample has a binary label: 1 = important, 0 = not important

**Example training matrix:**

$$X = \begin{bmatrix} 30 & 2 \\ 70 & 6 \\ 20 & 0 \\ 90 & 8 \end{bmatrix}, \quad y = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 1 \end{bmatrix}$$

## Python Code for Training

```
from sklearn.linear_model import LogisticRegression
import numpy as np

# Training data
X_train = np.array([
    [30, 2],
    [70, 6],
    [20, 0],
    [90, 8]
])
y_train = np.array([0, 1, 0, 1])

# Train model
model = LogisticRegression()
model.fit(X_train, y_train)
```

## Feature Extraction Example

```
def extract_features(article):
    title_length = len(article.get("title", ""))
    mentions = article.get("mention_count", 1)
    return [title_length, mentions]
```

## Prediction and Filtering

Once trained, the model can be used to predict the probability that new articles are important:

```
features = extract_features(article)
prob = model.predict_proba([features])[0][1]
if prob >= 0.7:
    mark_as_important(article)
```

## Why This Works

Logistic regression finds optimal weights for each feature, allowing us to build a probabilistic decision boundary. It's simple, fast, and interpretable — and can be retrained as new labeled data becomes available.

## 1.5 FDR-Based Selection (Benjamini–Hochberg Procedure)

This model uses the concept of **False Discovery Rate (FDR)** to control the number of false positives when selecting important articles. It is particularly useful when analyzing many items simultaneously — such as hundreds of news items — to avoid mistakenly selecting articles that appear important just by chance.

### False Discovery Rate (FDR)

FDR is defined as the expected proportion of false discoveries (type I errors) among the set of accepted hypotheses:

$$\text{FDR} = \mathbb{E} \left[ \frac{V}{R} \right]$$

Where:

- $V$  is the number of false positives,
- $R$  is the total number of rejected null hypotheses (selected articles).

Instead of controlling the probability of even a single false positive (as in Bonferroni), FDR allows some false positives but keeps their rate under control.

### Benjamini–Hochberg (BH) Procedure

The BH procedure operates as follows:

1. Sort the p-values  $p_1 \leq p_2 \leq \dots \leq p_n$ .
2. For each  $p_i$ , compute the threshold  $\frac{i}{n} \cdot \alpha$ , where  $\alpha$  is the desired FDR level.
3. Find the largest  $i$  such that  $p_i \leq \frac{i}{n} \cdot \alpha$ .
4. Reject all null hypotheses for  $p_1, \dots, p_i$ .

**Interpretation:** This guarantees that the expected proportion of false positives among the selected articles does not exceed  $\alpha$ .

### Pseudo P-values

In our unsupervised setting, we do not have real hypothesis tests. Instead, we simulate p-values inversely from article importance, using:

$$p_i = \frac{1}{\text{mention\_count}_i + \varepsilon}$$

Where  $\varepsilon$  is a small constant to prevent division by zero.

## Python Code

```
import numpy as np
from statsmodels.stats.multitest import multipletests

def compute_pvalues(articles):
    """
    Computes pseudo p-values for each article based on mention count.
    Lower p-value = more important.

    Parameters:
        articles (list): List of articles with 'mention_count'.

    Returns:
        list of float: Simulated p-values (1 / (mention_count + epsilon))
    """
    return [1.0 / (a.get("mention_count", 1) + 1e-8) for a in articles]

def select_important(articles, alpha=0.05):
    """
    Applies the Benjamini-Hochberg procedure to control the False Discovery Rate (FDR)
    when selecting significant articles.

    Parameters:
        articles (list): List of article dicts with 'mention_count'.
        alpha (float): Desired FDR level (default = 0.05).

    Returns:
        list: Filtered list of articles that pass FDR control.
    """
    if len(articles) == 0:
        return []

    # Compute pseudo p-values for each article
    pvals = compute_pvalues(articles)

    # Apply Benjamini-Hochberg FDR correction
    reject, _, _, _ = multipletests(pvals, alpha=alpha, method='fdr_bh')

    # Return only those articles that passed the FDR test
    return [a for a, keep in zip(articles, reject) if keep]
```

## Why This Works

The Benjamini-Hochberg method balances discovery and reliability. It allows us to select a wider range of potentially interesting articles while still controlling the false discovery rate — which is particularly valuable when monitoring large-scale, noisy news streams.

This is a statistically rigorous way to avoid overfitting to random spikes in mention counts.